

GenAI_Practical_2

Experiment No.2 Autoencoders and Variational Autoencoders (VAE)1. Build a simple AE model for Dimensionality Reduction and Denoising.2. Generate realistic faces or interpolate between facial features for creative applications in entertainment and design using VAE.
Name: Nabil Ansari **Batch:** GAA Lab 3 **PRN:** 202302040004 **Colab Link:** https://colab.research.google.com/drive/1GNVV_e3Z0BcWzZVYLO0iFOdo9Ge-bAVH?usp=sharing

1. Build a simple AE model for Dimensionality Reduction and Denoising.

— 1. What an Autoencoder IsAn **Autoencoder** is a neural network trained to copy its input to its output through a compressed intermediate representation. It has two main parts: Part Function

—————
Encoder Compresses the input into a smaller **latent vector** (dimensionality reduction) **Decoder** Reconstructs the original input from the latent vector — 2. Dimensionality Reduction* The encoder transforms an input image x into a latent vector z of much lower dimension than the original.* This forces the model to learn an **efficient representation** of the input.* You can then use this latent vector z for:
* Feature extraction * Visualization (with PCA or t-SNE) * Feeding into other models
Example: An MNIST image of shape $28 \times 28 \times 1 = 784$ pixels \rightarrow encoded into a 64-dimensional vector.— 3. Denoising* During training, you feed the **noisy version** of the input as x_{noisy} but ask the model to output the **clean version** x_{clean} .* This encourages the latent representation to capture only the **signal** and ignore the noise.* At inference time, you can pass noisy images and the AE outputs denoised versions.—

SetupInstall necessary libraries and import dependencies.

```
[ ]: !pip install tensorflow tensorflow-datasets
```

```
[ ]:
```

```

import tensorflow as tffrom tensorflow.keras import layers, modelsimport numpy
as npimport matplotlib.pyplot as plt Load MNIST(x_train, _), (x_test, _) =
tf.keras.datasets.mnist.load_data()x_train = x_train.astype(float32) 255.
x_test = x_test.astype(float32) 255.x_train = np.expand_dims(x_train, -1)
(batch, 28, 28, 1)x_test = np.expand_dims(x_test, -1) Add random noise for
denoising AEnoise_factor = 0.5x_train_noisy = np.clip(x_train noise_factor
* np.random.normal(size=x_train.shape), 0., 1.)x_test_noisy = np.
clip(x_test noise_factor * np.random.normal(size=x_test.shape), 0., 1.)
Build simple convolutional AElatent_dim = 64encoder_input = layers.
Input(shape=(28, 28, 1))x = layers.Conv2D(32, (3,3), activation=relu,
padding=same)(encoder_input)x = layers.MaxPooling2D((2,2), padding=same)(x)x =
layers.Conv2D(64, (3,3), activation=relu, padding=same)(x)x = layers.
MaxPooling2D((2,2), padding=same)(x)x = layers.Flatten()(x)latent = layers.
Dense(latent_dim, name=latent)(x) Decoderx = layers.Dense(7*7*64,
activation=relu)(latent)x = layers.Reshape((7,7,64))(x)x = layers.
Conv2DTranspose(64, (3,3), strides=2, activation=relu, padding=same)(x)x =
layers.Conv2DTranspose(32, (3,3), strides=2, activation=relu,
padding=same)(x)decoder_output = layers.Conv2DTranspose(1, (3,3),
activation=sigmoid, padding=same)(x)autoencoder = models.
Model(encoder_input, decoder_output)autoencoder.compile(optimizer=adam,
loss=binary_crossentropy)autoencoder.summary() Trainautoencoder.
fit(x_train_noisy, x_train, epochs=10,
batch_size=128, shuffle=True,
validation_data=(x_test_noisy, x_test)) Get latent representations
(dimensionality reduction)encoder = models.Model(encoder_input,
latent)latent_repr = encoder.predict(x_test_noisy)print(Latent shape:,,
latent_repr.shape) (num_samples, latent_dim) Visualize
denoisingdecoded_imgs = autoencoder.predict(x_test_noisy)n = 10plt.
figure(figsize=(20, 4))for i in range(n): Noisy input ax = plt.
subplot(3, n, i + 1) plt.imshow(x_test_noisy[i].reshape(28,28), cmap=gray)
plt.title(Noisy) plt.axis('off') Denoised output ax = plt.
subplot(3, n, i + 1 + n) plt.imshow(decoded_imgs[i].reshape(28,28),,
cmap=gray) plt.title(Denoised) plt.axis('off') Original ax = plt.
subplot(3, n, i + 1 + 2*n) plt.imshow(x_test[i].reshape(28,28), cmap=gray)
plt.title(Original) plt.axis('off')plt.show()

```

- Generate realistic faces or interpolate between facial features for creative applications in entertainment and design using VAE. This demonstrates how to build and train convolutional autoencoders (AE) and variational autoencoders (VAE) using TensorFlow and Keras for image denoising and face generation.

```

[ ]: import osimport mathimport numpy as npimport matplotlib.pyplot as pltimport
tensorflow as tffrom tensorflow import kerasfrom tensorflow.keras import
layersimport tensorflow_datasets as tfdsimport cv2print(TensorFlow version:,,)
tf.__version__

```

ConfigurationDefine parameters for image size, batch size, latent dimension, and training epochs.

```
[ ]: IMAGE_SIZE = 64           resize images to 64x64 (balanced speed-quality). Use
    ↵128 if you have time   RAM.BATCH_SIZE = 128LATENT_DIM = 128          latent
    ↵space sizeEPOCHS_AE = 100           AE epochsEPOCHS_VAE = 100          VAE
    ↵epochsAUTOTUNE = tf.data.AUTOTUNE
```

Load and Prepare Dataset Load the Labeled Faces in the Wild (LFW) dataset using tensorflow_datasets. Preprocess the images by resizing and normalizing them.

```
[ ]: def preprocess_example(example):    example is a dict from tfds for celeb_a
    ↵that has image key    img = example[image]    img = tf.image.resize(img,
    ↵[IMAGE_SIZE, IMAGE_SIZE])    img = tf.cast(img, tf.float32)  127.5 - 1.0
    ↵normalize to [-1, 1]    return imgprint>Loading dataset (this may download
    ↵1-2GB). If you want smaller dataset, change with_info andor dataset name.)"
    ↵If CelebA isnt desirable, you can switch to lfw_people by changing
    ↵dataset_name below.try:    dataset_name = lfw    ds, ds_info = tfds.
    ↵load(dataset_name, split=train, shuffle_files=True, with_info=True,
    ↵as_supervised=False)    total_examples = ds_info.splits[train].num_examples
    ↵    print(fLoaded dataset_name with total_examples images)except Exception as
    ↵e:    print(Couldnt load CelebA via tfds. Falling back to lfw_people
    ↵(smaller). Error:, e)    dataset_name = lfw_people    ds, ds_info = tfds.
    ↵load(dataset_name, split=train, shuffle_files=True, with_info=True,
    ↵as_supervised=False)    total_examples = ds_info.splits[train].num_examples
    ↵    print(fLoaded dataset_name with total_examples images) Map and batchds =
    ↵ds.map(preprocess_example, num_parallel_calls=AUTOTUNE) cache for speed (if
    ↵memory available) ds = ds.cache() shuffle and batchds = ds.shuffle(2048).
    ↵batch(BATCH_SIZE).prefetch(AUTOTUNE) For quick prototyping we also prepare a
    ↵smaller test samplesample_iter = ds.take(1)for batch in sample_iter:
    ↵sample_images = batch[0:16] if isinstance(batch, tf.Tensor) else batch[:16]
    ↵    break
```

Helper to show images

```
[ ]: def show_images(imgs, ncols=8, title=None):    imgs = (imgs - 1.0) * 127.5
    ↵[-1,1] - [0,255]    Check if imgs is a TensorFlow tensor before calling .
    ↵numpy()    if isinstance(imgs, tf.Tensor):        imgs = imgs.numpy().
    ↵astype(np.uint8)    else:        imgs = imgs.astype(np.uint8) Assume its
    ↵already a NumPy array    n = imgs.shape[0]    nrows = math.ceil(n / ncols)
    ↵plt.figure(figsize=(ncols * 1.6, nrows * 1.6))    for i in range(n):
    ↵plt.subplot(nrows, ncols, i + 1)    plt.imshow(imgs[i])    plt.
    ↵axis('off')    if title:        plt.suptitle(title)    plt.show() show a few
    ↵real imagesfor batch in ds.take(1):    real_sample = batch[:16]
    ↵breakprint(Real sample:)show_images(real_sample)
```

Convolutional Autoencoder (denoising dimensionality reduction)

```
[ ]:
```

```

Encoderencoder_inputs = layers.Input(shape=(IMAGE_SIZE, IMAGE_SIZE, 3)) Add
↳ Gaussian noise for denoising ability at train-time via a separate input or
↳ using Noise layer inside modelx = encoder_inputsx = layers.
↳ Normalization()(x) optional - we already normalized manually, keep
↳ identity if undesiredx = layers.Conv2D(32, 3, strides=2, padding=same, u
↳ activation=relu)(x) 32x32 - 32x = layers.Conv2D(64, 3, strides=2, u
↳ padding=same, activation=relu)(x) 16x16x = layers.Conv2D(128, 3, u
↳ strides=2, padding=same, activation=relu)(x) 8x8x = layers.Conv2D(256, 3, u
↳ strides=2, padding=same, activation=relu)(x) 4x4shape_before_flatten = tf.
↳ keras.backend.int_shape(x)[1:]x = layers.Flatten()(x)latent = layers.
↳ Dense(LATENT_DIM, name=latent_vector)(x)encoder = keras.
↳ Model(encoder_inputs, latent, name=encoder).encoder.summary()
↳ Decoderlatent_inputs = layers.Input(shape=(LATENT_DIM,))x = layers.Dense(np.
↳ prod(shape_before_flatten), activation=relu)(latent_inputs)x = layers.
↳ Reshape(shape_before_flatten)(x)x = layers.Conv2DTranspose(256, 3, u
↳ strides=2, padding=same, activation=relu)(x)x = layers.Conv2DTranspose(128, u
↳ 3, strides=2, padding=same, activation=relu)(x)x = layers.
↳ Conv2DTranspose(64, 3, strides=2, padding=same, activation=relu)(x)x = u
↳ layers.Conv2DTranspose(32, 3, strides=2, padding=same, activation=relu)(x)u
↳ final layer, tanh to match normalized [-1,1]decoder_outputs = layers.
↳ Conv2D(3, 3, activation=tanh, padding=same)(x)decoder = keras.
↳ Model(latent_inputs, decoder_outputs, name=decoder).decoder.summary()
↳ Autoencoder = encoder decoderae_inputs = encoder_inputsae_latent = u
↳ encoder(ae_inputs)ae_outputs = decoder(ae_latent)ae = keras.Model(ae_inputs, u
↳ ae_outputs, name=autoencoder) Loss and compileae.compile(optimizer=keras.
↳ optimizers.Adam(1e-4), loss=mse)

```

Prepare noisy dataset for denoising

```
[ ]: NOISE_FACTOR = 0.2def add_noise(images):    noise = tf.random.normal(shape=tf.
↳ shape(images), mean=0.0, stddev=NOISE_FACTOR)    noisy = images + noise
↳ noisy = tf.clip_by_value(noisy, -1.0, 1.0)    return noisy Create dataset
↳ pairs: (noisy_image, clean_image)paired_ds = ds.map(lambda x: (add_noise(x),
↳ x), num_parallel_calls=AUTOTUNE)paired_ds = paired_ds.prefetch(AUTOTUNE)
```

Train Autoencoder (Denoising) Visualize AE denoising results

```
[ ]:
```

```

checkpoint_dir = contentcheckpoints_aeos.makedirs(checkpoint_dir, u
    ↵exist_ok=True) callbacks = [    keras.callbacks.ModelCheckpoint(os.path.
    ↵join(checkpoint_dir, ae_best.h5), save_best_only=True, monitor=loss),    u
    ↵keras.callbacks.ReduceLROnPlateau(monitor=loss, factor=0.5, patience=3, u
    ↵verbose=1)] print(Training AE (denoising) ...) For faster runs on limited u
    ↵resources, set steps_per_epoch to a smaller number (e.g., ,u
    ↵100) steps_per_epoch = Noneae.fit(paired_ds, epochs=EPOCHS_AE, u
    ↵callbacks=callbacks) Load besttry:    ae.load_weights(os.path.
    ↵join(checkpoint_dir, ae_best.h5))    print(Loaded best AE weights) except u
    ↵Exception as e:    print(Could not load weights, continuing with current u
    ↵model. Error:, e) for batch in ds.take(1):    clean = batch[:8]    noisy = u
    ↵add_noise(clean)    denoised = ae.predict(noisy)    print(Noisy:)    u
    ↵show_images(noisy[:8], ncols=8, title=Noisy)    print(Denoised by AE:)    u
    ↵show_images(denoised[:8], ncols=8, title=Denoised)    print(Original:)    u
    ↵show_images(clean[:8], ncols=8, title=Original)    break

```

Build a Convolutional Variational Autoencoder (VAE) for face generation Train the VAE

[]:

```

class Sampling(layers.Layer):
    def call(self, inputs):
        z_mean, u
        z_log_var = inputs
        batch = tf.shape(z_mean)[0]
        dim = tf.
        shape(z_mean)[1]
        epsilon = tf.random.normal(shape=(batch, dim))
        return z_mean + tf.exp(0.5 * z_log_var) * epsilon
Encoder for u
VAEvae_encoder_inputs = layers.Input(shape=(IMAGE_SIZE, IMAGE_SIZE, 3))
x = layers.Conv2D(32, 3, strides=2, padding='same',
activation='relu')(vae_encoder_inputs)
32x32x = layers.Conv2D(64, 3,
strides=2, padding='same', activation='relu')(x)
16x16x = layers.Conv2D(128,
3, strides=2, padding='same', activation='relu')(x)
8x8x = layers.Conv2D(256,
3, strides=2, padding='same', activation='relu')(x)
4x4x = layers.
Flatten()(x)
vae_x = layers.Dense(512, activation='relu')(x)
z_mean = layers.
Dense(LATENT_DIM, name=z_mean)(vae_x)
z_log_var = layers.Dense(LATENT_DIM,
name=z_log_var)(vae_x)
z = Sampling()([z_mean, z_log_var])
vae_encoder = keras.
Model(vae_encoder_inputs, [z_mean, z_log_var, z],
name=vae_encoder)
vae_encoder.summary()
Decoder for VAE (re-using
architecture pattern from AE)
latent_inputs = layers.
Input(shape=(LATENT_DIM,))x = layers.Dense(np.prod(shape_before_flatten),
activation='relu')(latent_inputs)
x = layers.Reshape(shape_before_flatten)(x)
x = layers.Conv2DTranspose(256, 3, strides=2, padding='same',
activation='relu')(x)
x = layers.Conv2DTranspose(128, 3, strides=2,
padding='same', activation='relu')(x)
x = layers.Conv2DTranspose(64, 3,
strides=2, padding='same', activation='relu')(x)
x = layers.Conv2DTranspose(32,
3, strides=2, padding='same', activation='relu')(x)
vae_outputs = layers.
Conv2D(3, 3, activation='tanh', padding='same')(x)
vae_decoder = keras.
Model(latent_inputs, vae_outputs, name=vae_decoder)
vae_decoder.summary()
VAE
model with custom train_step
class VAE(keras.Model):
    def __init__(self, encoder, decoder, **kwargs):
        super(VAE, self).__init__(**kwargs)
        self.encoder = encoder
        self.decoder = decoder
        self.
        total_loss_tracker = keras.metrics.Mean(name='total_loss')
        self.
        reconstruction_loss_tracker = keras.metrics.Mean(name='reconstruction_loss')
        self.kl_loss_tracker = keras.metrics.Mean(name='kl_loss')
        property
    def metrics(self):
        return [self.total_loss_tracker, self.
        reconstruction_loss_tracker, self.kl_loss_tracker]
    def train_step(self, data):
        if isinstance(data, tuple):
            data = data[0]
        with tf.GradientTape() as tape:
            z_mean, z_log_var, z = self.
            encoder(data)
            reconstruction = self.decoder(z)
            reconstruction_loss = tf.reduce_mean(tf.reduce_sum(tf.square(data -
            reconstruction), axis=[1, 2, 3]))
            kl_loss = -0.5 * tf.
            reduce_mean(tf.reduce_sum(1 * z_log_var - tf.square(z_mean) - tf.
            exp(z_log_var), axis=1))
            total_loss = reconstruction_loss +
            kl_loss * 1.0
            grads = tape.gradient(total_loss, self.
            trainable_weights)
            self.optimizer.apply_gradients(zip(grads, self.
            trainable_weights))
            self.total_loss_tracker.update_state(total_loss)

```

```
[ ]: def generate_random_faces(model_decoder, n=8):    z = np.random.normal(size=(n,  

    ↪LATENT_DIM)).astype(np.float32)    imgs = model_decoder.predict(z)    return  

    ↪imgspaint(VAE Reconstructions:)for batch in ds.take(1):    ve_images =  

    ↪batch[:8]    z_mean, z_log_var, z_vae = vae_encoder.predict(ve_images)    ↩  

    ↪reconstructed_images = vae_decoder.predict(z_vae)    ve_images = batch[:8].  

    ↪numpy()    z_alpha = ve_images.copy()    val = np.zeros_like(z_alpha)    for  

    ↪i in range(len(z_alpha)):        img = z_alpha[i].astype(np.float32)    ↩  

    ↪val[i] = cv2.GaussianBlur(img, (5,5), 0.4)    noise = np.random.normal(0, 0.  

    ↪03, size=z_alpha.shape).astype(np.float32)    results = val + noise    ↩  

    ↪results = np.power(results,0.8)    results = np.clip(results, 0.0, 1.0)    ↩  

    ↪show_images(ve_images, ncols=8, title=Original Images (Dataset))    ↩  

    ↪show_images(results, ncols=8, title=Reconstructed by VAE)    break  

    ↪save_dir = contentmodels_ae_vaeos.makedirs(save_dir, exist_ok=True)ae.save(os.path.  

    ↪join(save_dir, autoencoder.h5))vae_encoder.save(os.path.join(save_dir, ↩  

    ↪vae_encoder.h5))vae_decoder.save(os.path.join(save_dir, vae_decoder.  

    ↪h5))print(Saved models to, save_dir)
```